

External Validation of a Seven-Weight Scorecard versus Tree Ensembles in Ovarian Cancer Risk Prediction

Jiayi Yao
BASIS Independent Bellevue
Bellevue, WA, USA
yaojiayi2020@gmail.com

Zooley Lane Go Hua
Issaquah High School
Issaquah, WA, USA
Zooley.hua@gmail.com

Vritika S Sharma
Issaquah High School
Issaquah, WA, USA
vritikas01@gmail.com

Abstract—Preoperative blood tests help decide whether an ovarian mass requires cancer surgery; a model that degrades silently at a new hospital can delay surgery or trigger unnecessary operations. This paper asks whether more complex machine learning models transfer better to new hospitals than simple ones. Eight models (a CA125 cutoff rule, the clinical formulas ROMA and CPH-I, a logistic regression, an integer scorecard, and three tree ensembles) were trained on a Chinese dataset (349 patients) and applied unchanged elsewhere. External evidence is deliberately narrow: the primary benchmark is 100 West China patients with measured biomarkers (63 cancers), reported alongside the same cohort with HE4 imputed for 74 percent (380 patients); a Japanese cohort (177 patients, 150 healthy women) serves only as a specificity stress test, not as an external validation cohort. On the 100 complete cases, the scorecard (area under the ROC curve, 0.971) was numerically highest and did not differ significantly from the logistic regression (0.941) or ROMA (0.965). A computer simulation, calibrated to the real data and intended only as an illustrative counterexample, found that the flexible models won when the biomarker-cancer relationship stayed constant, but under the simulated form of concept drift examined here their advantage diminished and reversed (the scorecard was most accurate in 22 of 24 conditions). Within these small Asian cohorts, a hand-computable integer scorecard performed no worse than far more complex tree ensembles and missed fewer cancers than the off-the-shelf formulas; whether these findings generalize beyond these cohorts requires larger validation studies.

Index Terms—ovarian cancer; risk prediction; scorecard; external validation; transportability; concept drift

I. Introduction

Ovarian cancer is the most lethal gynecologic cancer, and whether a pelvic mass goes to a specialist cancer surgeon depends on preoperative risk; getting it wrong in either direction is costly. That risk is assessed from two blood proteins measured cheaply before surgery, CA125 (cancer antigen 125) and HE4 (human epididymis protein 4), combined either by fixed formulas such as the Risk of Ovarian Malignancy Algorithm (ROMA) [1] and the Copenhagen Index (CPH-I) [2] or by a fitted model. Machine learning models proposed for this task usually show strong training results but little external testing [3], [4].

This paper asks whether model complexity buys performance that survives a move to a new hospital, and

whether a bedside-computable scorecard matches the complex alternatives. We trained eight models on one Chinese cohort, applied them unchanged to a second Chinese cohort and a Japanese specificity stress test, and measured discrimination, calibration, and decision consequences, plus a simulation exploring when complex models help.

II. Methods

A. Data

Three public, anonymized datasets were used [6]–[8]. Institutional review board approval was not required for re-analysis of public, de-identified data; the cohorts are a training cohort from Changzhou, Jiangsu, China (349 patients; 171 ovarian cancers, 178 benign masses; Mendeley Data, doi:10.17632/th7fztbrv9.11 [9]), a second Chinese cohort from the West China Second University Hospital of Sichuan University, Chengdu (380 patients; 188 cancers, 192 benign cysts; figshare, doi:10.6084/m9.figshare.28831256), and a Japanese cohort from a university hospital in Japan (177 patients; 27 cancers, 150 healthy women; Karger figshare, doi:10.6084/m9.figshare.24235450, [8]). Each patient has four features (CA125 U/mL, HE4 pmol/L, age, menopausal status). The raw files are chinese training dataset(All Raw Data).csv, west_china.xlsx, and japanese.xlsx; src/clean_data.py harmonizes them into the three processed files used by every analysis. In the training cohort, scattered missing biomarker values were filled by feature-wise cohort medians; externally, HE4 was not measured for 74% of West China patients and age was not recorded for 85% of Japan patients; these external gaps were filled with the training-set medians, computed once and frozen, so that no statistic is ever computed on a validation cohort. The external validation evidence is therefore deliberately narrow. The primary West China analyses use the 100 patients with measured HE4 (63 cancers, 37 benign). Complete-case analysis is primary for three reasons: (i) the alternative, training-median imputation, replaces 74% of HE4 values with one constant, which cannot reflect between-patient variation

and artificially narrows uncertainty; (ii) multiple imputation by chained equations (MICE, Section III-C) degrades the scorecard’s tertile binning (0.853), so imputed benchmarks depend on the imputation method; and (iii) the complete cases evaluate the exact inputs a clinician would measure. The 280 excluded patients (125 cancers, 155 benign) had similar age, CA125, and menopausal rates. The 74%-imputed full cohort ($N = 380$) is reported alongside as a co-primary analysis. Because it contains 150 healthy women and only 27 cancers, the Japanese cohort is treated as a specificity stress test rather than an external validation of the triage tool: its AUROC is reported separately in Table II and enters the pooled analysis only as a discrimination sensitivity check, never to support pooled decision or clinical claims.

B. Models

Eight models used the same four features. (1) Integer scorecard: each continuous feature was rank-transformed to a uniform score (QuantileTransformer, 100 quantile points) and cut into training-set tertiles, with frozen cut-points CA125 24.3/165.0 U/mL, HE4 45.8/92.7 pmol/L, and age 37.0/52.0 years; the three bins per feature were one-hot encoded and fitted with an L2-regularized logistic regression ($C = 1.5$). Each coefficient β_i was converted to an integer weight $w_i = \text{round}(5\beta_i/\max_j |\beta_j|) \in \{-5, \dots, +5\}$, with $|\beta_i| < 0.01$ set to zero; the scaling factor was $5/\max_j |\beta_j| = 2.151$, so the rounding is exactly reproducible. The deployed score sums the integer weights of the patient’s bins (range -7 to $+9$, computable by hand; Table I) and uses seven nonzero weights (HE4 $-4/-1/+5$, CA125 $-1/+2$, age $-2/+2$); the CA125 and age mid-bins round to zero, and the fitted intercept is not used in scoring. The binary menopausal feature, binned identically, collapses to one constant column whose coefficient rounds to zero; it is effectively dropped, not manually removed (its row in Table I documents this). (2) Logistic regression on standardized features. (3–4) the published formulas ROMA and CPH-I, computed on raw values. ROMA uses $\text{PI} = -12.0 + 2.38 \ln(\text{HE4}) + 0.0626 \ln(\text{CA125})$ premenopause and $\text{PI} = -8.09 + 1.04 \ln(\text{HE4}) + 0.732 \ln(\text{CA125})$ postmenopause, mapped to a percentage by the logistic transform, with published cutoffs 13.1% (pre) and 27.7% (post) [1]. CPH-I uses $-14.0647 + 1.0649 \log_2(\text{HE4}) + 0.6050 \log_2(\text{CA125}) + 0.2672 \times \text{age}/10$, likewise logistic-transformed, with cutoff 7% [2]. Units: HE4 in pmol/L, CA125 in U/mL, age in years; menopause is the dataset’s binary postmenopausal indicator; imputed HE4 uses the frozen training median. Within-cohort recalibration means fitting, on West China, a logistic regression of the outcome on the formula’s probability as the sole covariate. (5) a single CA125 ≥ 35 U/mL rule. (6–8) XGBoost, CatBoost, and Random Forest, evaluated under two settings: Setting A keeps each package’s default hyperparameters, and Setting B tunes each ensemble by internal cross-validation over a

TABLE I
The integer scorecard, fitted on the Chinese training cohort (349 patients) and applied unchanged externally; features binned by frozen training-set tertiles.

Feature	Low bin	Mid bin	High bin
CA125 (U/mL)	−1	0	+2
HE4 (pmol/L)	−4	−1	+5
Age (years)	−2	0	+2
Menopausal status	0	0	0

27-configuration grid per family, with all tuning decisions confined to training folds and the external cohorts never entering selection (XGBoost 100 trees, depth 3, learning rate 0.03; CatBoost 500 iterations, depth 4, learning rate 0.03; Random Forest 500 trees, depth 6, minimum leaf size 2). Setting B is the primary comparison (Table II; Setting A in the table footnote). All models were trained once on the Chinese cohort; every parameter, bin cutpoint, and threshold was frozen before external evaluation.

C. Evaluation

Discrimination was measured by the area under the ROC curve (AUROC). Confidence intervals (95%) and two-sided paired tests of between-model differences used 2,000 patient-level bootstrap resamples, which preserve the paired structure across models and keep Monte Carlo error below 0.001 AUROC. The prespecified primary comparison is the scorecard versus ROMA, tested per cohort and summarized by a DerSimonian–Laird random-effects model [5] across the imputed West China cohort and the Japanese stress-test cohort; with only two cohorts the heterogeneity estimate is imprecise, so per-cohort paired tests are the primary evidence and the pooled estimate is a summary. All other comparisons are secondary, and effect sizes with 95% CIs are emphasized over p-values. Calibration was measured with the Brier score and calibration-in-the-large. Decision value was quantified with a weighted-harm analysis (Section III-B): decision thresholds were swept over 1–100%, and for exchange rates $w \in \{1, 5, 10, 50\}$ the threshold minimizing $H = w \times \text{missed} + \text{unnecessary}$ per 100 women was identified. We distinguish Q1 (does each model transport unchanged?) from Q2 (how well can each model perform after local recalibration?). The exchange-rate weights are hypothetical scenarios, not validated clinical utilities, spanning equal valuation of a missed cancer and an unnecessary referral ($w = 1$) to valuing a missed cancer 50 times as much ($w = 50$); they characterize sensitivity to relative error costs. All design choices preceded external evaluation: the four features are those shared by all three cohorts and required by the published formulas (the primary comparator is ROMA, the most widely used published formula); the primary comparison (scorecard versus ROMA) and the four weights were pre-specified; only the scorecard’s bin count and regularization were selected among 12 pre-specified candidates

(bins 3/4/5/7 with $C = 0.5/1.5/5.0$) by repeated CV (Section III-E). The primary harm comparison applies within-cohort recalibration to the scorecard and the two formulas (the logistic regression and XGBoost are shown as fitted); published-cutpoint operating points are secondary, and no pooled decision analysis was performed. The retraining experiment used stratified training subsets of size $n \in \{60, 90, 120, 180, 240, 349\}$ (10 draws each), evaluated unchanged on the imputed West China cohort. Robustness was tested with 10% and 30% multiplicative, additive, and combined noise on CA125 and HE4, 5% gross-error contamination, and $\pm 10\%$ bin-boundary shifts.

D. Implementation and availability

Analyses used Python 3.13 with scikit-learn 1.7.2, XGBoost 3.3.0, CatBoost 1.2.10, and Random Forest, all with random seed 42; cross-validation used stratified five-fold splits. Setting A kept the package defaults (XGBoost 100 trees, depth 6, learning rate 0.3; CatBoost 1,000 iterations, depth 6, learning rate 0.03; Random Forest 100 trees, no depth cap); Setting B used the internal-CV-selected configurations of Section II-B. The complete code and results are available at <https://github.com/angelayao/Ovarian-Cancer-RiskPrediction>, with a one-command pipeline (reproduce.sh) and pinned dependencies.

E. Simulation

The simulation is an illustrative counterexample, not an explanation of the observed real-world results, and no general claim about real concept drift is made from it. A simulation calibrated to the Chinese data drew biomarkers from class-conditional lognormal distributions and generated the cancer label from the logistic model

$$\eta(z_1, z_2, z_3) = -0.65 + 1.8z_1 + 1.0z_2 + 0.5z_3 + (1-d)(1.6z_1z_2 + 1.2\sin(1.7z_1)), \quad (1)$$

where z_1 , z_2 , and z_3 are standardized log CA125, log HE4, and age, and $d \in [0, 1]$ sets the drift: $d = 0$ keeps the interaction and nonlinear terms, $d = 1$ removes them. The grid crossed training sizes $n \in \{100, 200, 400, 800\}$, distribution shifts $s \in \{0, 0.5, 1.0\}$, measurement noise $\{0\%, 50\%\}$, and drift levels $d \in \{0, 1\}$, with 20 replicates per cell.

III. Results

A. Discrimination

On the 100 West China patients with measured HE4 (63 cancers; the primary external benchmark), the scorecard was numerically highest, with ROMA, tuned XGBoost, and the logistic regression within overlapping CIs (Table II, West CC column; Fig. 1).

On the co-primary 74%-imputed cohort ($N = 380$) the ranking was similar (logistic regression 0.936, scorecard 0.929, tuned ensembles 0.877–0.895; Table II, West column).

TABLE II
AUROC (95% CI) by model and cohort.

Model	West CC	West	Japan
Scorecard	0.971 0.936–0.995 0.941	0.929 0.902–0.955 0.936	0.845 0.714–0.949 0.806
Logistic reg.	0.887–0.983 0.965	0.908–0.961 0.886	0.666–0.926 0.799
ROMA	0.930–0.992 0.929	0.848–0.921 0.868	0.666–0.908 0.920
CPH-I	0.875–0.973 0.694	0.828–0.905 0.642	0.852–0.973 0.781
CA125 rule	0.605–0.785 0.946	0.597–0.684 0.890	0.685–0.874 0.840
XGBoost	0.897–0.984 0.910	0.853–0.925 0.877	0.711–0.944 0.772
CatBoost	0.845–0.962 0.930	0.836–0.915 0.895	0.650–0.880 0.819
Random Forest	0.871–0.979	0.858–0.931	0.692–0.926

CC: complete case; CI: confidence interval. West CC ($n=100$): measured-HE4 complete cases (primary benchmark); West ($N=380$): the 74%-imputed cohort (co-primary); Japan ($n=177$): 150 healthy women and 27 cancers (specificity stress test). Tree ensembles are Setting B (internally CV-tuned); Setting A (defaults) scored 0.921/0.922/0.910 (CC), 0.887/0.879/0.879 (West), 0.736/0.750/0.794 (Japan).

As a discrimination benchmark, a random-effects (DerSimonian–Laird) meta-analysis with paired bootstrap tests pooling the imputed West China cohort ($N = 380$) and the Japanese stress-test cohort ($n = 177$) gave a pooled scorecard-minus-ROMA difference of +0.043 AUROC (95% CI 0.013 to 0.074; $p = 0.006$). The scorecard also exceeded each tuned tree ensemble (pooled differences +0.036, +0.055, and +0.033 for XGBoost, CatBoost, and Random Forest; all $p \leq 0.016$) and did not differ from the logistic regression (-0.005 ; $p = 0.55$). After Holm–Bonferroni correction across the five prespecified pooled comparisons (scorecard versus ROMA, the logistic regression, and the three tuned tree ensembles), the scorecard–ROMA difference remained significant (adjusted $p = 0.023$), as did all three ensemble comparisons (adjusted p at most 0.031). These pooled comparisons combine the imputed cohort with a stress-test cohort and are secondary to the complete-case benchmark above.

In the Japan specificity stress test ($n = 177$, of whom only 27 have cancer), CPH-I was highest (0.920; Table II, Japan column), with wide confidence intervals given so few cancers.

B. Decision consequences

In the primary weighted-harm analysis (Q2; Table III, Fig. 2), the scorecard and the two formulas were within-cohort recalibrated (logistic regression and XGBoost as fitted). The scorecard’s minimum harm was below recalibrated ROMA and CPH-I at weights up to 5 (12.1 vs. 14.7 and 15.0 at $w = 1$); at $w = 50$ both fell below it (45.8 and 50.3 versus 59.5). The logistic regression was lowest at weights up to 10; at $w = 50$, recalibrated




Fig. 1. External validation AUROC, 100 West China patients with measured HE4 (63 cancers); dots show bootstrap means and 95% CIs (2,000 resamples).




Fig. 2. Decision curve analysis on West China ($N = 380$): net benefit against the decision threshold, with treat-all (dashed) and treat-none (horizontal axis) references; the shaded band marks the range of published cutpoints (7–30%). Scorecard and formulas are shown after within-cohort recalibration (logistic regression as fitted); XGBoost omitted.

ROMA was lowest (45.8). Reported minima are over the swept thresholds; where a minimum exceeds the treat-all reference (e.g., XGBoost at $w = 50$), treating everyone dominates that model (reference row in Table III).

As the Q1 decision consequence, the scorecard’s high-risk tier (score $\geq +1$) missed 47 of 188 cancers with 5 false positives, versus 82 (ROMA) and 48 (CPH-I) at the published cutpoints.

C. Calibration and missing-data checks

For calibration on the imputed West China cohort ($N = 380$), the scorecard’s logistic mapping was refit on that same cohort and evaluated there, so its calibration figures

TABLE III
Minimum net harm on West China ($N = 380$) at exchange-rate weight w : harm = $w \times \text{missed} + \text{unnecessary}$ per 100 women, minimized over the swept thresholds. Q2 rows: scorecard, ROMA, and CPH-I are within-cohort recalibrated; the logistic regression and XGBoost are shown as fitted. Q1 published-scale harms: ROMA 15.0/43.4/45.8/45.8; CPH-I 15.0/51.3/64.7/159.5.

Model	$w=1$	$w=5$	$w=10$	$w=50$
Logistic reg.	9.7	29.2	39.7	46.6
Scorecard	12.1	30.5	45.0	59.5
ROMA (recalibrated)	14.7	42.9	45.8	45.8
CPH-I (recalibrated)	15.0	50.3	50.3	50.3
XGBoost	14.7	44.7	66.1	215.5
Treat all (reference)	50.5	50.5	50.5	50.5

are apparent, not prospective; the logistic regression is shown frozen from training and the formulas at published scales, so only the scorecard’s figures can be inflated by the shared fit-and-evaluate data. The recalibrated scorecard had Brier 0.097, calibration slope 1.00, and expected calibration error (ECE) 0.042; the logistic regression had Brier 0.111, slope 1.44, ECE 0.135; ROMA had Brier 0.209, slope 1.53, ECE 0.220; CPH-I had Brier 0.259, slope 1.11, ECE 0.303, with both formulas over-predicting risk (calibration-in-the-large -0.95 and -1.42).

Multiple imputation by chained equations (MICE) preserved the flexible-model ranking (logistic regression 0.926) but degraded the scorecard’s tertile binning (0.853; see Section II-A).

D. Robustness and simulation

Under 30% multiplicative noise on CA125 and HE4 (West China, $N = 380$; 100 perturbations), mean AUROC losses were 0.039 (scorecard), 0.002 (logistic regression), 0.073 (ROMA), and 0.084 (XGBoost).

In the simulation (Fig. 3), at 50% noise the ensemble-average minus-scorecard gap averaged $+0.018$ at drift level $d = 0$ of (1) and -0.018 at $d = 1$; at $d = 1$ the scorecard had the highest AUROC in 22 of 24 cells. Across factors, the gap spread most across drift levels (0.033), then distribution shift (0.012), training size (0.003), and noise (0.002).

E. Internal validity and sensitivity

Across 20 repeated five-fold cross-validations, the scorecard’s internal AUROC was 0.906 (SD 0.003) and the three-bin configuration was selected in 18 of 20 repeats; in-sample ensembles fitted at 0.991–1.000, 0.094–0.117 below their CV AUROC.

IV. Discussion and Conclusion

The results show three things, within the limits of the available cohorts. First, on discrimination: seven integer weights performed no worse than models with thousands of parameters; the scorecard was numerically best on the 100 complete cases (0.971), similar to the logistic regression on the imputed cohort (0.929 vs. 0.936), and CPH-I won




Fig. 3. Regime map from the simulation: mean over 20 replicates of the tree-ensemble-average minus scorecard test AUROC at 50% noise (red favors the ensembles, blue the scorecard); axes: training samples (100–800), distribution shift (0–1).

the Japan stress test (0.920). Second, on decisions: the published cutpoints missed 82 of 188 cancers versus 47 for the locally fitted scorecard, whose minimum harm stayed below the recalibrated formulas up to 5:1. Third, under the simulated form of concept drift examined here, the flexible models’ advantage diminished and reversed; the simulation does not by itself show that this reversal occurs in real hospitals.

The main limitations are as follows. (i) The external samples are small: the primary complete-case cohort contains 100 patients (63 cancers), the Japanese cohort 27 cancers, and all three cohorts are Asian hospital series; these analyses reduce but do not remove this uncertainty, and the conclusions may not generalize beyond these cohorts. (ii) HE4 was missing for 74% of West China patients and age for 85% of Japan patients, so much of the external comparison rests on training-set constants; and the Japanese cohort consists of healthy volunteers rather than women with masses, so its discrimination is not directly comparable with West China’s; both narrow how broadly the findings apply, although complete-case and imputed results agreed. (iii) Ultrasound-based models could not be tested (the public data lack imaging), and the simulation is a simplification, not a proof about any specific hospital.

Future work should add ultrasound variables and validate prospectively in larger, multi-center cohorts; on these small cohorts the scorecard already performed no worse than far more complex models and missed fewer cancers than the off-the-shelf formulas.

Acknowledgment

The authors acknowledge the use of Google Gemini, an AI-based assistant, for grammar and language checks and for basic code debugging; the authors retain full authority and responsibility for the design, methodology, analysis, and all claims of this paper.

References

- [1] R. G. Moore, M. C. Miller, P. Disilvestro, et al., “Evaluation of the diagnostic accuracy of the risk of ovarian malignancy algorithm in women with a pelvic mass,” *Obstet. Gynecol.*, vol. 118, pp. 280–288, 2011.
- [2] M. A. Karlsen, E. V. S. Høgdall, I. J. Christensen, et al., “A novel diagnostic index combining HE4, CA125 and age may improve triage of women with suspected ovarian cancer,” *Gynecol. Oncol.*, vol. 138, pp. 640–646, 2015.
- [3] E. Kawakami, J. Tabata, N. Yanaihara, et al., “Application of artificial intelligence for preoperative diagnostic and prognostic prediction in epithelial ovarian cancer based on blood biomarkers,” *Clin. Cancer Res.*, vol. 25, pp. 3006–3015, 2019.
- [4] E. Christodoulou, J. Ma, G. S. Collins, E. W. Steyerberg, J. Y. Verbakel, and B. Van Calster, “A systematic review shows no performance benefit of machine learning over logistic regression for clinical prediction models,” *J. Clin. Epidemiol.*, vol. 110, pp. 12–22, 2019.
- [5] R. DerSimonian and N. Laird, “Meta-analysis in clinical trials,” *Control. Clin. Trials*, vol. 7, no. 3, pp. 177–188, 1986.
- [6] Q. Mi, J. Jiang, T. Znati, et al., “Data for: Using machine learning to predict ovarian cancer,” *Mendeley Data*, vol. 11, 2020, doi:10.17632/th7fzbrv9.11.
- [7] Li, “Comparative analysis of clinical characteristics between ovarian cancer and ovarian cyst patients,” *figshare*, 2025, doi:10.6084/m9.figshare.28831256.
- [8] H. Machida, T. Hirakawa, K. Tsunekawa, et al., “Revised cut-off value of human epididymis protein 4 enhances its use as an ovarian tumor marker,” *Gynecol. Obstet. Invest.*, vol. 88, pp. 349–358, 2023.
- [9] M. Lu, Z. Fan, B. Xu, et al., “Using machine learning to predict ovarian cancer,” *Int. J. Med. Inform.*, vol. 141, p. 104195, 2020.